

Network Planner Performance

Optimus: Machine where we run "big" NP scenarios
32G RAM, Dual Quad Core, SSD (which didn't survive)
and yet...

Num Nodes	Grid Size	Memory Consumption	Run Time
937	3696	3G	339 minutes
5577	4641	36G	3966 minutes

That last scenario originally took even longer due to excessive paging.

Run Time

Problem:

Profiling revealed intersection test as the dominant time consumer.
(Up to 99% as scenario size increases)

For each candidate segment that passes mvMax criteria, check intersection with every segment in network.

If 20k segments pass, and there are 1k segments in network: 20M intersection tests. (each is several steps)

Solution:

Introduce R-Tree to store/search network for intersections.

Now for each candidate segment, there are $\log_b(K)$ intersection calls.

b is tree branching factor and K is the number segments in the network.

So for the same scenario above, if branching factor is 2, there are $\frac{1}{10}$ th of the intersection tests.

Memory

Problem:

There are $\frac{(n)(n-1)}{2} = \binom{n}{2}$ candidate segments (n is the number of nodes in the scenario)

For 5k nodes, that's $> 12\text{M}$ candidate segments

Each candidate segment is a Python object, originally consuming 1k, for a total of 12G

Solution:

Represent candidate segments as a 3xN matrix

uint16, uint16, float32 \rightarrow node index 1, node index 2, length
8 bytes each

So now the candidate segments for the scenario above consume 96M

Results

Before:

Num Nodes	Grid Size	Memory Consumption	Run Time
937	3696	3G	339 minutes
5577	4641	36G	3966 minutes

After:

Num Nodes	Grid Size	Memory Consumption	Run Time
937	3696	1G	27 minutes
5577	4641	1.2G	276 minutes

Practically, 10X memory and run-time reduction for “large” scenarios

Up to 60X memory reduction within modKruskal

More dramatic improvements as scenario size grows